

https://colab.research.google.com/github/google-gemini/cookbook/blob/main/templates/aistudio_gemini_prompt_freeform.ipynb

```
# import necessary modules.
import base64
import copy
import json
import pathlib
import requests

import PIL.Image
import IPython.display
from IPython.display import Markdown

try:
    # The SDK will automatically read it from the GOOGLE_API_KEY
    environment variable.
    # In Colab get the key from Colab-secrets ("🔑" in the left panel).
    import os
    from google.colab import userdata

    os.environ["GOOGLE_API_KEY"] = userdata.get("GOOGLE_API_KEY")
except ImportError:
    pass

import google.generativeai as genai

# Parse the arguments

model = 'gemini-2.0-flash-exp' # @param {isTemplate: true}
contents_b64 =
'W3sicGFydHMiOlt7InRleHQiOiJpbXBvcnQgdG9yY2hcXG5pbXBvcnQgdG9yY2gubm4gYXMgb
m5cXG5pbXBvcnQgdG9yY2gub3B0aW0gYXMgb3B0aW1cXG5mcm9tIHRvcnNodmlzaW9uIGltcG9
ydCBkYXRhc2V0cywgdHJhbnNmb3Jtc1xcbmZyb20gdG9yY2gudXRpbHMuZGF0YSBpbXBvcnQgR
GF0YUxvYWRlc1xcbmZyb20gc2tsZWYybi5tZXRYaWNzIGltcG9ydCBjbGFzc2lmaWNhdGlvbl9
yZXBvcnRcXG5pbXBvcnQgbWF0cGxvdGxpYi5weXBsb3QgYXMgcGx0XFxuXFxuIyBEZXZpY2UgY
29uZmlndXJhdGlvblxcbmRldmljZSA9IHRvcnNoLmRldmljZSgnY3VkYScgaWYgdG9yY2guY3V
kYS5pc19hdmFpbGFibGUoKSB1bHN1ICdjchUNKVxcblxcbiMgSHlwZXJwYXJhbWV0ZXJzXFxub
nVtX2Vwb2NocyA9IDNcXG5iYXRjaF9zaXplID0gNjRcXG5sZWYybmluZ19yYXRlID0gMC4wMVx'
```

[illegible]

gQ05OTW9kZWwoKS50byhkZXZpY2UpXFxuY3JpdGVyaW9uID0gbm4uQ3Jvc3NFbnRyb3B5TG9zc
ygpXFxub3B0aWlpemVyID0gb3B0aW0uQWRhbShtb2RlbC5wYXJhbWV0ZXJzKCksIGxyPWxlYXJ
uaW5nX3JhdGUpXFxuc2NoZWR1bGVyID0gb3B0aW0ubHJfc2NoZWR1bGVyLlN0ZXBMUihvcHRpb
Wl6ZXIsIHN0ZXBfc2l6ZT0xLCBnYWltYT0wLjcpXFxuXFxuIyBUcmFpbmluZyBmdW5jdGlvblx
cbmRlZiB0cmFpbihb2RlbCwgZGV2aWNlLCB0cmFpb19sb2FkZXIsIG9wdGltaxPlciwgY3Jpd
GVyaW9uKTpcXG4gICAgbW9kZWwudHJhaW4oKVxcbiAgICBlcG9jaF9sb3NzID0gMFxcbiAgICB
mb3IgYmF0Y2hfaWR4LCAoZGF0YSwgdGFyZ2V0KSBpbilB1bnVtZXJhdGUodHJhaW5fbG9hZGVyK
TpcXG4gICAgICAgIGRhGEsIHRhcmdldCA9IGRhGEudG8oZGV2aWNlKSwgdGFyZ2V0LnRvKGR
ldmljZS1cXG4gICAgICAgIG9wdGltaxPlci56ZXJvX2dyYWQoKVxcbiAgICAgICAgb3V0cHV0I
D0gbW9kZWwoZGF0YS1cXG4gICAgICAgIGxvc3MgPSBjcml0ZXJpb24ob3V0cHV0LCB0YXJnZXQ
pXFxuICAgICAgICBsb3NzLmJhY2t3YXJkKClcXG4gICAgICAgIG9wdGltaxPlci5zdGVwKClcX
G4gICAgICAgIGVwb2NoX2xvc3MgKz0gbG9zcy5pdGVtKClcXG4gICAgcmV0dXJuIGVwb2NoX2x
vc3MgLyBsZW4odHJhaW5fbG9hZGVyKVxcblxcbiMgVGVzdGluZyBmdW5jdGlvblxcbmRlZiB0Z
XN0KG1vZGVsLCBkZXZpY2UsIHRlc3RfbG9hZGVyLCBjcml0ZXJpb24pOlxcbiAgICBtb2RlbC5
ldmFsKClcXG4gICAgdGVzdF9sb3NzID0gMFxcbiAgICBjb3JyZWNOID0gMFxcbiAgICBhbGxfc
HJlZHMgPSBbXVxcbiAgICBhbGxfdGFyZ2V0cyA9IFtdXFxuICAgIHdpdGggdG9yY2gubm9fZ3J
hZCgpOlxcbiAgICAgICAgZm9yIGRhGEsIHRhcmdldCBpbilB0ZXN0X2xvYWRlcjpcXG4gICAgI
CAgICAgICBkYXRhLCB0YXJnZXQgPSBkYXRhLnRvKGRldmljZSksIHRhcmdldC50byhkZXZpY2U
pXFxuICAgICAgICAgICAgb3V0cHV0ID0gbW9kZWwoZGF0YS1cXG4gICAgICAgICAgICB0ZXN0X
2xvc3MgKz0gY3JpdGVyaW9uKG9ldHBldCwgdGFyZ2V0KS5pdGVtKCKgKiBkYXRhLnNpemUoMC1
cXG4gICAgICAgICAgICBwcmVklD0gb3V0cHV0LmFyZ21heChkaW09MSwga2VlcGRpbT1UcnVlK
VxcbiAgICAgICAgICAgIGNvcnJlY3QgKz0gcHJlZC5lcSh0YXJnZXQudmlld19hcyhwcmVlKSk
uc3VtKCKuaXRlbSgpXFxuICAgICAgICAgICAgYWxsX3ByZWZzLmV4dGVuZChwcmVlLmNwdSgpL
m51bXB5KCKpXFxuICAgICAgICAgICAgYWxsX3RhcmdldHMuZXh0ZW5kKHRhcmdldC5jcHUoKS5
udW1weSgpKVxcblxcbiAgICB0ZXN0X2xvc3MgLz0gbGVuKHRlc3RfbG9hZGVyLmRhZGFzZXQpX
FxuICAgIGFjY3VyYWN5ID0gMTAwLiAqIGNvcnJlY3QgLyBsZW4odGVzdF9sb2FkZXIuZGF0YXN
ldC1cXG4gICAgcHJpbmQoZidcXFxcblRlc3Qgc2V0OiBBdmVyYWdlIGxvc3M6IHT0ZXN0X2xvc
3M6LjRmfSwgQWNjdXJhY3k6IHTjb3JyZWNOfs97bGVuKHRlc3RfbG9hZGVyLmRhZGFzZXQpfSA
oe2FjY3VyYWN5Oi4wZn0lKVxcXFxuJylcXG5cXG4gICAgcmV0dXJuIHRlc3RfbG9zcywgYWNjd
XJhY3ksIGFsbf9wcmVlcYwgYWxsX3RhcmdldHNcXG5cXG4jIFRyYWluaW5nIGxvb3Agd2l0aCB
2aXNlYWxpemF0aW9uXFxudHJhaW5fbG9zc2VzID0gW11cXG50ZXN0X2xvc3NlcY9IFtdXFxuY
WNjdXJhY2llcyA9IFtdXFxuXFxuZm9yIGVwb2NoIGluIHJhbmdlKDEsIG51bV9lcG9jaHMgKyA
xKTpcXG4gICAgcHJpbmQoZlxcXCJFcG9jaCB7ZXBvY2h9L3tudW1fZXBvY2hzfVxcXCJpXFxuI
CAgIHRyYWluaX2xvc3MgPSB0cmFpbihb2RlbCwgZGV2aWNlLCB0cmFpb19sb2FkZXIsIG9wdGl
taxPlciwgY3JpdGVyaW9uKVxcbiAgICB0ZXN0X2xvc3MsIGFjY3VyYWN5LCBhbGxfcHJlZHMzI
GFsbF90YXJnZXRzID0gdGVzdChtb2RlbCwgZGV2aWNlLCB0ZXN0X2xvYWRlcjpcXG4gICAgI
HJhY3ksIGFsbf9wcmVlcYwgYWxsX3RhcmdldHMuZXh0ZW5kKGFjY3VyYWN5K5Vx
cbiAgICBzY2hlZHVzZXIuc3RlcCgpXFxuXFxuIyBDbGFzc2lmaWNhdGlvbiByZXBvcnRcXG5w
cludChcXFwiXFxcXG5DbGFzc2lmaWNhdGlvbiBSZXBvcnQ6XFxcIilcXG5wcmVudChjbGFzc2l

```

maWNhdGlvbl9yZXBvcnQoYWxsX3RhcmdldHMsIGFsbF9wcmVkeywgdGFyZ2V0X25hbWVzPXRyY
WluX2RhdGFzZXQuY2xhc3NlcypkXFxuXFxuIyBQbG90IHRyYWluaW5nIGxvc3NcXG5wbHQuZml
ndXJlKGZpZ3NpemU9KDEwLCA1KS1cXG5wbHQucGxvdCh0cmFpb19sb3NzZXMsIGxhYmVsPSdUc
mFpbmluZyBMb3NzJylcXG5wbHQucGxvdCh0ZXN0X2xvc3NlcywgbGFIZWw9J1Rlc3QgTG9zeYc
pXFxucGx0LnhsYWJlbCgnRXBvY2hzJylcXG5wbHQucGxvdCh0ZXN0X2xvc3NlcywgbGFIZWw9J1Rlc3QgTG9zeYc
GUoJ1RyYWluaW5nIGFuZCBUZXR0IEExvc3MnKVxcbnBsdC5sZWdlbmQoKVxcbnBsdC5zaG93KC1
cXG5cXG4jIFNhdmUgdGhlIG1vZGVsXFxudG9yY2guc2F2ZShtb2RlbCwgJ2Zhc2hpb25fbW5pc
3RfY25uX2Z1bGwucHRoJylcXG5wcm1udChcXFwiTW9kZWwgc2F2ZWQgYXMgJ2Zhc2hpb25fbW5
pc3RfY25uX2Z1bGwucHRoJy5cXFwiKSJ9LHsidGV4dCI6Im1ucHV0OiAifSx7InRleHQiOiJvd
XRwdXQ6ICJ9XX1d' # @param {isTemplate: true}

generation_config_b64 =
'eyJ0ZWlwZXJhdHVyZSI6MSwidG9wX3AiOjAuOTUsInRvcF9rIjo0MCwibWF4X291dHBldF90b
2t1bnMiOjgOTJ9' # @param {isTemplate: true}

safety_settings_b64 = "e30=" # @param {isTemplate: true}

gais_contents = json.loads(base64.b64decode(contents_b64))

generation_config = json.loads(base64.b64decode(generation_config_b64))
safety_settings = json.loads(base64.b64decode(safety_settings_b64))

stream = False

# Convert and upload the files

tempfiles = pathlib.Path(f"tempfiles")
tempfiles.mkdir(parents=True, exist_ok=True)

drive = None
def upload_file_data(file_data, index):
    """Upload files to the Files API.

    For each file, Google AI Studio either sent:
    - a Google Drive ID,
    - a URL,
    - a file path, or
    - The raw bytes (`inline_data`).

    The API only understands `inline_data` or it's Files API.

```


This code, uploads files to the files API where the API can access them.

```
"""

mime_type = file_data["mime_type"]
if drive_id := file_data.pop("drive_id", None):
    if drive is None:
        from google.colab import drive
        drive.mount("/gdrive")

    path = next(
pathlib.Path(f"/gdrive/.shortcut-targets-by-id/{drive_id}").glob("*")
    )
    print("Uploading:", str(path))
    file_info = genai.upload_file(path=path, mime_type=mime_type)
    file_data["file_uri"] = file_info.uri
    return

if url := file_data.pop("url", None):
    response = requests.get(url)
    data = response.content
    name = url.split("/")[-1]
    path = tempfiles / str(index)
    path.write_bytes(data)
    print("Uploading:", url)
    file_info = genai.upload_file(path, display_name=name,
mime_type=mime_type)
    file_data["file_uri"] = file_info.uri
    return

if name := file_data.get("filename", None):
    if not pathlib.Path(name).exists():
        raise IOError(
            f"local file: `{name}` does not exist. You can upload
files "
            'to Colab using the file manager (" Files" in the left '
            "toolbar)"
        )
)
```

```
        file_info = genai.upload_file(path, display_name=name,
mime_type=mime_type)
        file_data["file_uri"] = file_info.uri
        return

    if "inline_data" in file_data:
        return

    raise ValueError("Either `drive_id`, `url` or `inline_data` must be
provided.")

contents = copy.deepcopy(gais_contents)

index = 0
for content in contents:
    for n, part in enumerate(content["parts"]):
        if file_data := part.get("file_data", None):
            upload_file_data(file_data, index)
            index += 1

import json
print(json.dumps(contents, indent=4))
```



```
import torch
import torch.nn as nn
import torch.nn.functional as F # Import F
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt

# Device configuration
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Hyperparameters
num_epochs = 7
batch_size = 64
learning_rate = 0.03

# Data augmentation and normalization
transform = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)) # Normalize to [-1, 1]
])

# Load the Fashion-MNIST Dataset
train_dataset = datasets.FashionMNIST(root='./data',
                                       train=True,
                                       transform=transform,
                                       download=True)

test_dataset = datasets.FashionMNIST(root='./data',
                                       train=False,
                                       transform=transforms.Compose([
                                           transforms.ToTensor(),
                                           transforms.Normalize((0.5,),
(0.5,))
                                       ]),
                                       download=True)
```



```

# Data loaders
train_loader = DataLoader(dataset=train_dataset,
                           batch_size=batch_size,
                           shuffle=True)

test_loader = DataLoader(dataset=test_dataset,
                          batch_size=batch_size,
                          shuffle=False)

# Define the CNN Model
class CNNModel(nn.Module):
    def __init__(self):
        super(CNNModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 8, 5)
        self.bn1 = nn.BatchNorm2d(8) # Batch Normalization
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(8, 16, 5)
        self.bn2 = nn.BatchNorm2d(16) # Batch Normalization
        self.dropout = nn.Dropout(0.25) # Dropout layer
        self.fc1 = nn.Linear(16 * 4 * 4, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.bn1(self.conv1(x))))
        x = self.pool(F.relu(self.bn2(self.conv2(x))))
        x = x.view(-1, 16 * 4 * 4)
        x = F.relu(self.fc1(x))
        x = self.dropout(x) # Apply dropout
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

# Initialize the model, loss function, and optimizer
model = CNNModel().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=1, gamma=0.7)

# Training function

```

```

def train(model, device, train_loader, optimizer, criterion):
    model.train()
    epoch_loss = 0
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = criterion(output, target)
        loss.backward()
        optimizer.step()
        epoch_loss += loss.item()
    return epoch_loss / len(train_loader)

# Testing function
def test(model, device, test_loader, criterion):
    model.eval()
    test_loss = 0
    correct = 0
    all_preds = []
    all_targets = []
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += criterion(output, target).item() * data.size(0)
            pred = output.argmax(dim=1, keepdim=True)
            correct += pred.eq(target.view_as(pred)).sum().item()
            all_preds.extend(pred.cpu().numpy())
            all_targets.extend(target.cpu().numpy())

    test_loss /= len(test_loader.dataset)
    accuracy = 100. * correct / len(test_loader.dataset)
    print(f'\nTest set: Average loss: {test_loss:.4f}, Accuracy:
{correct}/{len(test_loader.dataset)} ({accuracy:.0f}%) \n')

    return test_loss, accuracy, all_preds, all_targets

# Training loop with visualization
train_losses = []
test_losses = []

```

```

accuracies = []

for epoch in range(1, num_epochs + 1):
    print(f"Epoch {epoch}/{num_epochs}")
    train_loss = train(model, device, train_loader, optimizer, criterion)
    test_loss, accuracy, all_preds, all_targets = test(model, device,
test_loader, criterion)
    train_losses.append(train_loss)
    test_losses.append(test_loss)
    accuracies.append(accuracy)
    scheduler.step()

# Classification report
print("\nClassification Report:")
print(classification_report(all_targets, all_preds,
target_names=train_dataset.classes))

# Plot training loss
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Training Loss')
plt.plot(test_losses, label='Test Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Test Loss')
plt.legend()
plt.show()

# Save the model
torch.save(model, 'fashion_mnist_cnn_full.pth')
print("Model saved as 'fashion_mnist_cnn_full.pth'.")

```

Epoch 1/3

Test set: Average loss: 0.4532, Accuracy: 8378/10000 (84%)

Epoch 2/3


Test set: Average loss: 0.3807, Accuracy: 8594/10000 (86%)

Epoch 3/3

Test set: Average loss: 0.3399, Accuracy: 8760/10000 (88%)

Classification Report:

	precision	recall	f1-score	support
T-shirt/top	0.78	0.87	0.82	1000
Trouser	1.00	0.97	0.98	1000
Pullover	0.76	0.87	0.81	1000
Dress	0.87	0.91	0.89	1000
Coat	0.80	0.79	0.80	1000
Sandal	0.97	0.96	0.97	1000
Shirt	0.73	0.52	0.61	1000
Sneaker	0.92	0.96	0.94	1000
Bag	0.96	0.98	0.97	1000
Ankle boot	0.96	0.94	0.95	1000
accuracy		0.88		10000
macro avg	0.87	0.88	0.87	10000
weighted avg	0.87	0.88	0.87	10000



Model saved as 'fashion_mnist_cnn_full.pth'.

Epoch 1/5

Test set: Average loss: 0.5421, Accuracy: 8116/10000 (81%)

Epoch 2/5

Test set: Average loss: 0.4704, Accuracy: 8319/10000 (83%)

Epoch 3/5

Test set: Average loss: 0.4342, Accuracy: 8404/10000 (84%)

Epoch 4/5


Test set: Average loss: 0.4015, Accuracy: 8504/10000 (85%)

Epoch 5/5

Test set: Average loss: 0.4019, Accuracy: 8487/10000 (85%)

Classification Report:

	precision	recall	f1-score	support
T-shirt/top	0.72	0.88	0.79	1000
Trouser	0.97	0.98	0.98	1000
Pullover	0.64	0.90	0.75	1000
Dress	0.91	0.86	0.88	1000
Coat	0.79	0.68	0.73	1000
Sandal	0.98	0.92	0.95	1000
Shirt	0.70	0.38	0.49	1000
Sneaker	0.90	0.98	0.94	1000
Bag	0.96	0.97	0.96	1000
Ankle boot	0.97	0.94	0.95	1000
accuracy		0.85		10000
macro avg	0.85	0.85	0.84	10000
weighted avg	0.85	0.85	0.84	10000



Model saved as 'fashion_mnist_cnn_full.pth'.

Epoch 1/7

Test set: Average loss: 0.5426, Accuracy: 8065/10000 (81%)

Epoch 2/7

Test set: Average loss: 0.5100, Accuracy: 8144/10000 (81%)

Epoch 3/7

Test set: Average loss: 0.4346, Accuracy: 8403/10000 (84%)

Epoch 4/7

Test set: Average loss: 0.4262, Accuracy: 8425/10000 (84%)

Epoch 5/7

Test set: Average loss: 0.4010, Accuracy: 8543/10000 (85%)

Epoch 6/7


Test set: Average loss: 0.3877, Accuracy: 8581/10000 (86%)

Epoch 7/7

Test set: Average loss: 0.3773, Accuracy: 8603/10000 (86%)

Classification Report:

	precision	recall	f1-score	support
T-shirt/top	0.78	0.81	0.79	1000
Trouser	0.99	0.97	0.98	1000
Pullover	0.77	0.77	0.77	1000
Dress	0.85	0.90	0.88	1000
Coat	0.74	0.82	0.78	1000
Sandal	0.97	0.92	0.94	1000
Shirt	0.68	0.54	0.60	1000
Sneaker	0.90	0.96	0.93	1000
Bag	0.95	0.97	0.96	1000
Ankle boot	0.95	0.95	0.95	1000
accuracy			0.86	10000
macro avg	0.86	0.86	0.86	10000
weighted avg	0.86	0.86	0.86	10000



Model saved as 'fashion_mnist_cnn_full.pth'.